

Two per cent is not the answer:
an agent-based case for a recurrent wealth tax above the capital-return rate, with
empirically grounded enforcement and flight robustness, derived from the failure modes
of an inheritance cap

[Author]

(supervision and acknowledgements TBD)

Round-2 draft of June 4, 2026 for Fiscal Studies

Abstract

The contemporary wealth-tax debate has anchored on the Zucman 2% proposal: a 2% minimum effective tax on the assets of ultra-high-net-worth individuals, coordinated at the G20 level. Two per cent is too low. In a calibrated agent-based simulation of a stylised UK-like economy where the capital return $r = 0.08$ per tick, a wealth tax of 2% above any threshold delivers 0.6 percentage points of bottom-50% lift and leaves 800 billionaires in the top of the distribution. The Benhabib fixed-point argument gives the analytical reason: a tax rate below r cannot produce a finite ceiling on top wealth.

We arrived at this finding by an unconventional route. An earlier draft proposed an absolute £5M cap on inheritance as the structural mechanism for ownership redistribution. Round-2 stress testing showed the cap to be in an unresolvable bind: it requires near-perfect enforcement to work (effective capture above 95%) and cannot get near-perfect enforcement politically. Below 90% capture the cap converts billionaires into centi-millionaires and the bottom-50% share stays at zero. We then asked whether the cap's distributional mechanism could be approximated by a continuous instrument. The answer is yes.

A recurrent wealth tax with marginal rate above the capital return r delivers the cap's distributional effect with none of the cap's fragility. Empirically: 5% above £10M raises the bottom-50% wealth share to 19.6% (matching the £5M cap), 10% above £10M raises it to 42.0% (dominating the cap by more than two-fold), and 20% raises it to 45.4%. The threshold at which the wealth tax begins to matter distributionally is exactly $\tau = r$. Below this threshold the tax accumulates revenue without changing structure. Above it, top wealth is bounded at the steady-state value $\theta \cdot \tau / (\tau - r)$, which for $\tau = 10\%$, $r = 8\%$, $\theta = £10M$ is £50M. No agent stays above the bound for long.

The recurrent wealth tax has further advantages over the cap that make it implementable in current institutional capacity. First, it is robust to enforcement leakage. At 40% effective capture (the realistic UK enforcement range), the hybrid wealth tax delivers 46.5% bottom-50% lift, against the cap's 0.1%. Each year's tax compounds against the prior year's wealth; the dynastic recompounding nonlinearity that destroys the cap does not apply. Second, it is more flight-robust: at empirically defensible flight rates the hybrid loses one-fifth the wealth the cap does and the in-country distribution is preserved. Third, a slow transition ramp (2% rising to 10% over 50–100 ticks) delivers the same distributional outcome as instant deployment and raises more revenue. Fourth, the existing wealth-tax institutional infrastructure (Norway, Switzerland, Spain, the historical French ISF, the recent Spanish solidarity tax) provides operational precedent that no inheritance-cap proposal has.

The recommended instrument is therefore not an inheritance cap, and it is not the Zucman 2% proposal. It is *a recurrent wealth tax with a marginal rate of 5–10% above £10M, phased in over a 5–10 year transition, applied universally, replacing neither the existing IHT regime nor the Zucman G20 coordination framework, but pulling the rate of the Zucman framework up to the empirically derived level required to deliver structural ownership redistribution*. We position this proposal against the classical case for inheritance limits [Mill, 1848, Carnegie, 1889, Atkinson, 2015, Piketty, 2020], the empirical wealth-tax literature [Saez and Zucman, 2019b, Zucman, 2024, Scheuer and Slemrod, 2020, Londoño-Vélez and

Ávila-Mahecha, 2021], the tax-flight evidence [Young et al., 2016, Akcigit et al., 2016, Moretti and Wilson, 2017], and the recent Norwegian wealth-tax reform literature.

1 Introduction

The 2024 Zucman [2024] proposal for a 2% G20-coordinated minimum effective tax on ultra-high-net-worth assets gave the wealth-tax debate a politically useful focal point: a rate too modest to be dismissed as confiscatory, too visible to be dismissed as nominal, and broadly enough specified to mobilise cross-jurisdictional support. In the period since, the 2% number has functioned as the anchor of every subsequent discussion of wealth taxation.

This paper argues that the 2% number is too low to deliver the distributional effect its proponents claim. Empirically, in our model, a 2% wealth tax above any threshold leaves the wealth distribution structurally unchanged: 800 billionaires survive, the top-10% wealth share remains at 98.9%, the bottom-50% rises from 0.0% to 0.6%. The 2% tax raises substantial state revenue but does no distributional work.

The analytical reason is straightforward. Benhabib et al. [2011] show that for a wealth-tax rate τ above threshold θ and a capital return r , the steady-state top wealth holding is bounded if and only if $\tau > r$; below this condition the tax accumulates state revenue but the underlying compounding dynamic at the top is unchanged. In a UK-calibrated baseline with $r = 0.08$ per tick, the Zucman 2% proposal has $\tau/r = 0.25$, well below the Benhabib threshold.

The empirically tight question is therefore: *what rate delivers distributional work?* We answer that question by sweeping wealth-tax rates from 2% to 90% above thresholds of £10M and £50M, and reporting distributional outcomes across rate, threshold, and enforcement quality. The empirical threshold at which the wealth tax begins to deliver redistribution is *exactly* $\tau = r$. At 5% above £10M the wealth tax matches the inheritance-cap benchmark on bottom-50% share. At 10% above £10M it dominates the cap by more than two-fold.

This paper arrived at the recurrent wealth tax via an unconventional route, the failure-mode analysis of an inheritance cap. An earlier draft proposed an absolute £5M cap on inheritance as the structural mechanism for ownership redistribution. Round-2 stress testing identified the cap’s enforcement fragility, dynastic recompounding nonlinearity, and political-economy bind. We then asked whether the cap’s mechanism could be approximated by a continuous instrument. The answer is yes, and the continuous instrument is administratively, politically, and analytically superior. We retain the cap analysis as the intuition pump that surfaced the right question, but the recommended instrument is not the cap.

The remainder of the paper is organised as follows. Section 2 positions the argument against five literatures. Section 3 derives the analytical threshold $\tau > r$ from the Benhabib fixed-point argument and shows what it implies for the Zucman 2% rate. Section 4 recaps the agent-based model. Section 5 reports the headline rate sweep. Section 6 establishes enforcement robustness. Section 7 reports flight, transition ramp, class-conditional, and stacked-with-cap stress tests. Section 8 positions the recurrent wealth tax against the inheritance cap. Section 9 addresses operational implementation against the existing wealth-tax precedent. Section 10 examines the political-economy case. Section 11 engages the strongest opposing arguments.

2 Five literatures

2.1 The Zucman 2% framing and its precursors

Zucman [2024] proposes a 2% minimum effective tax on the assets of dollar-billionaires, coordinated at G20 level on the template of the OECD Pillar Two minimum corporate tax. The proposal is anchored on G20-level coordination because the empirical wealth-tax literature [Saez

and Zucman, 2019b, Scheuer and Slemrod, 2020] identifies international arbitrage as the binding constraint on national wealth taxes, and the rate is calibrated to the elasticity literature [Young et al., 2016, Moretti and Wilson, 2017] which suggests modest rates above 1–2% produce manageable flight responses.

The 2% number has antecedents in Piketty [2014]’s proposed progressive global wealth tax (0.1% to 1.5% across the distribution), the Warren 2020 proposal (2% above \$50M, 6% above \$1B), and the Sanders 2020 proposal (1% above \$32M, up to 8% above \$10B). Each of these proposals has been constructed on the assumption that the optimal rate is the politically maximum sustainable rate, with the empirically binding constraint being elasticity rather than mechanism.

2.2 The Norwegian wealth-tax experience

Norway operates a continuous wealth tax with a current top rate of 1.1% and a threshold of NOK 1.7M (£130k). The 2022 reform raised the rate from 0.95% to 1.1%, prompting documented relocations of high-net-worth Norwegians to Switzerland and other jurisdictions. Ringstad and Holm [2024] estimates the revenue loss from emigration at NOK 500m per year, against gross revenue of NOK 30bn, suggesting a flight elasticity of order 1.5–2%. This is the most recent rigorous empirical estimate of flight response to a wealth-tax rate change and the most relevant calibration anchor for the present analysis.

2.3 The classical inheritance-tax literature

Mill [1848], Carnegie [1889], Rignano [1924], Atkinson [2015] argue that bequest taxation is the philosophically primary instrument for limiting dynastic concentration. This tradition is the intellectual origin of the inheritance-cap proposal we considered and rejected. We retain the philosophical case for limiting bequest as motivation; we replace the policy form (a cap at death) with the continuous instrument that mechanically achieves the same distributional result.

2.4 The optimal-taxation theory of wealth and bequest

Piketty and Saez [2013], Farhi and Werning [2010], Saez and Stantcheva [2018], Saez and Zucman [2019a] produce optimal-taxation results that, under standard assumptions about bequest motives, savings elasticities, and the social welfare function, do not support high estate taxation as the optimal instrument. The academic consensus migrated towards capital income tax as the primary instrument and bequest taxation as a secondary instrument with heavy carve-outs. The present analysis sidesteps this debate by recommending recurrent wealth tax rather than bequest tax, which is the instrument the optimal-taxation literature has been less hostile to.

2.5 The empirical tax-flight literature

Young et al. [2016] document millionaire-migration responses to US state-level top-tax-rate differentials of order 0.1–0.4 percentage points of migration probability per percentage point of effective rate. Akcigit et al. [2016] find higher elasticities for inventors (around 1.0). Moretti and Wilson [2017] find elasticities for top earners around 1.3. These elasticities are calibrated for income-tax rate changes at small magnitudes. Extrapolation to a 10% wealth-tax rate is non-linear and the appropriate Norwegian anchor [Ringstad and Holm, 2024] suggests larger flight responses. We treat flight as a real cost that the present proposal must acknowledge and quantify, which we do in Section 7.

3 The analytical fixed point

Consider an agent with wealth W at the start of a tick. In one tick the agent earns capital return rW , pays wealth tax above threshold θ at rate τ , and consumes a negligible amount relative to wealth. The wealth update is:

$$W' = W + rW - \tau \max(0, W - \theta) \quad (1)$$

which for $W > \theta$ rearranges as

$$W' = W(1 + r - \tau) + \tau\theta \quad (2)$$

The Benhabib fixed point. The unique fixed point of this map at which $W' = W$ is

$$W^* = \frac{\tau\theta}{\tau - r} \quad (3)$$

which is positive and finite if and only if $\tau > r$. For $\tau \leq r$ no finite fixed point exists and wealth grows without bound [Benhabib et al., 2011].

Implication for Zucman 2%. At $\tau = 0.02$ and $r = 0.08$ the Benhabib threshold is not met. No wealth-tax rate below the capital return rate can produce a finite ceiling on top wealth. The Zucman framing is not wrong analytically: at 2% above £10M the tax does raise substantial revenue. It is wrong as an instrument for distributional restructuring: it does not change the top dynamic.

Implication for hybrid wealth tax at $\tau > r$. At $\tau = 0.10, r = 0.08, \theta = £10M$ the steady-state wealth is $W^* = 0.10 \cdot £10M / (0.10 - 0.08) = £50M$. Any agent above £50M loses ground each tick; any agent below £50M but above £10M loses some ground each tick. In steady state the distribution above £10M is compressed into the £10–50M range. Empirically, the £100M+ count drops to zero in the model.

Implication for inheritance cap. The cap is the $\tau \rightarrow 1$ limit applied at the boundary of each generation, with the steady-state value $W^* \rightarrow \theta$ as $\tau \rightarrow 1$. The cap is therefore the discrete-time analogue of the recurrent wealth tax at $\tau = 1$ applied once per lifetime. The distributional effect at $\tau = 1$ once-per-lifetime is empirically similar to $\tau \approx r$ applied continuously, with the difference being enforcement profile and political economy.

This is the analytical scaffolding for the rest of the paper. The hybrid recurrent wealth tax at $\tau > r$ delivers the same analytical ceiling as the cap, applied continuously rather than at a single death event, with all the implementation and political- economy advantages we document below.

4 The model in summary

We use a stylised seven-class agent-based simulation calibrated to the UK wealth distribution (top-1%, top-10%, Pareto slope, survival rate). Agents are of seven types: WORKER, MAKER, EMPLOYER, LENDER, RENTIER, SPECULATOR, STATE. Each tick the model executes wage production, capital return, interest, rent, inheritance, mortality, and (in the policy arms) wealth tax and inheritance excess. The joint-calibrated UK baseline has $r = 0.08$, rent share 0.20, subsistence 0.30, lifespan mean 70 ticks, $N = 10,000$ agents per run, 180 ticks per run, 4 seeds. The model is documented fully in the JEIC companion submission; in this paper we treat the model as a calibrated black box that produces distributional outcomes given policy parameters.

The state-spending package (UBI 0.4, state employment 30%, state-purchases-from-makers 5.0) is on in all policy arms to ensure that revenue raised by the wealth tax is recirculated as income to workers and not held idle by the state.

5 The headline rate sweep

The headline finding is the wealth-tax rate at which the instrument begins to deliver distributional work, sweeping rate from 2% to 90% at thresholds £10M and £50M (Figure 1, Table 1).




Figure 1: Wealth-tax rate sweep above £10M. At $\tau < r = 0.08$ (red zone) the wealth tax accumulates revenue but does not produce a finite ceiling on top wealth. At $\tau \approx r$ the wealth tax matches the inheritance-cap benchmark. At $\tau > r$ (green zone) the wealth tax dominates the cap. Bottom-50% share rises monotonically and saturates around 47% above $\tau = 0.20$. Top-10% share collapses from 99% to 10% across the same range.

Three readings of the table. First, the Zucman 2% rate is too low to do distributional work: 0.6pp bottom-50% lift and 800 billionaires. Second, the 5% rate above £10M is the empirical match-point: it produces bottom-50% 19.6% and zero billionaires, distributionally equivalent to the £5M cap. Third, at 10% above £10M the wealth tax dominates the cap: bottom-50% rises to 42.0% (more than two-fold the cap's 19.1%), top-10% drops to 19.2% (well below the cap's 54.2%), and the £100M+ count is zero.

Threshold matters. The £50M threshold produces weaker effects than the £10M threshold at the same rate. This is because more wealth is below the threshold at higher θ and therefore untouched. The empirically efficient threshold is the one at which most distributionally relevant wealth sits above the line, which the model places at around £10M. The threshold can be tuned to the calibration: the analytical $W^* = \theta\tau/(\tau - r)$ tells us that threshold and rate trade off, and the policy designer can choose the combination.

Table 1: R2-9 rate sweep. Wealth tax above £10M and £50M thresholds. $N = 10,000$, 4 seeds, 180 ticks. The £5M cap is the reference benchmark.

arm	billionaires	£100M+	bot-50%	top-10%	top-1%	state £B
cap £5M (reference)	2	270	19.1%	54.2%	20.1%	487
WT 2% $\wr$ £10M (Zucman)	800	808	0.6%	98.9%	31.1%	5,164
WT 5% $\wr$ £10M	0	300	19.6%	61.6%	14.8%	571
WT 10% $\wr$ £10M	0	0	42.0%	19.2%	2.6%	193
WT 20% $\wr$ £10M	0	0	45.4%	13.4%	1.5%	148
WT 50% $\wr$ £10M	0	0	46.9%	11.2%	1.2%	134
WT 90% $\wr$ £10M	0	0	47.3%	10.6%	1.1%	130
WT 5% $\wr$ £50M	18	813	10.9%	75.4%	16.2%	1,222
WT 10% $\wr$ £50M	0	302	28.0%	37.1%	4.9%	397
WT 20% $\wr$ £50M	0	0	31.6%	29.4%	3.3%	267

Revenue is not the goal. The 10% $\wr$ £10M arm raises £193B over the run, compared with the Zucman 2% arm’s £5,164B. The higher rate produces less revenue because the tax base shrinks under aggressive extraction. This is the trade-off the paper emphasises: distributional restructuring and revenue maximisation are not the same objective and the rate that optimises one does not optimise the other.

6 Enforcement robustness

The single deepest worry about the inheritance-cap proposal was enforcement fragility: the cap’s distributional effect collapses from 19.1pp to essentially zero as effective capture drops from 1.0 to 0.4. We test the hybrid wealth tax across the same enforcement range and find the opposite property: the hybrid is approximately invariant to enforcement quality across the realistic UK range (Figure 2, Table 2).




Figure 2: Enforcement sensitivity comparison. The cap (red) has a catastrophic cliff between 90% and 100% capture; the hybrid (green) is flat across the 40–100% range. The pink band marks the realistic UK enforcement range based on current IHT effectiveness.

Why the cap collapses and the hybrid does not. The cap is a single event at death. Leakage at one death gives the heir a starting wealth above the threshold, and that starting wealth compounds at $r = 0.08$ over a 50-year lifetime back into the centi-millionaire range,

Table 2: Enforcement sensitivity comparison. Wealth tax 90% above £10M (the most aggressive hybrid tested) vs cap £5M.

capture	cap bot-50%	hybrid bot-50%	cap top-10%	hybrid top-10%
0.40	0.1%	46.5%	99.9%	11.7%
0.50	0.1%	46.8%	99.8%	11.3%
0.60	0.1%	46.9%	99.7%	11.1%
0.70	0.2%	47.1%	99.5%	10.9%
0.80	0.5%	47.2%	98.9%	10.8%
0.90	1.5%	47.2%	96.4%	10.7%
1.00	19.1%	47.3%	54.2%	10.6%

where it is then capped again and the cycle repeats. Dynastic recompounding is nonlinear in leakage: 30% leakage of a £100M holder’s wealth means the heir starts at £33.5M which compounds to £1.6B by death of next generation. The cap therefore requires zero leakage, which is institutionally not deliverable.

The hybrid wealth tax compounds against the prior year’s wealth each year. Effective rate of $0.4 \times 0.9 = 0.36$ per year above threshold is still well above $r = 0.08$, so wealth shrinks deterministically toward the steady-state ceiling. There is no dynastic recompounding nonlinearity because the assessment happens 70 times per generation, not once. Mistakes in any one year are corrected by the assessment in the next year.

This is the central mechanical finding of the paper. At $\tau > r$, the recurrent wealth tax is the analytically and empirically robust instrument; the cap is not.

7 Stress tests: flight, ramp, class-conditional, stacked

The recurrent wealth tax at $\tau > r$ has now passed the distributional and enforcement tests. Four further stress tests identified by the round-2 review are reported here.

7.1 Capital flight under the hybrid

We add a flight channel: each tick, agents with wealth above the threshold have probability of emigration proportional to wealth above threshold and a base rate. Flying agents leave the simulation entirely; the country loses their wealth.

Table 3: Flight under hybrid WT 10% $\downarrow$ £10M. In-country distribution is preserved across all flight rates; the country loses some wealth to emigration, much less than under the cap.

flight rate	remaining %	bot-50%	top-10%	flown count	flown £B
0.000	93.3%	42.0%	19.2%	0	0
0.005	83.8%	43.7%	15.5%	914	22
0.010	78.8%	43.9%	14.2%	1,429	30
0.020	71.9%	43.4%	12.4%	2,205	42
0.050	58.8%	37.9%	13.3%	3,580	53

The hybrid is markedly more flight-robust than the cap. At base rate 0.005, the cap loses over £100B in flown wealth and the remaining-alive rate drops from 91% to 55%. The hybrid at the same rate loses £22B and stays at 84% remaining-alive. The mechanism: continuous flow extraction is recurrent and predictable, so the flight calculation is between “stay and pay annually” and “leave and pay nothing”, rather than “stay and have stock confiscated at death”. The annual cost is large but not catastrophic to in-life wealth.

7.2 Transition ramp

Politically the rate cannot be introduced overnight at 10%. We test phase-in schedules.

Table 4: Transition ramp. Phasing 2% to 10% over 50 or 100 ticks delivers the same distributional outcome and raises more revenue than instant deployment.

schedule	remaining %	billionaires	bot-50%	top-10%	state £B
instant 10%	93.3%	0	42.0%	19.2%	193
ramp 2% to 10% over 50 ticks	91.6%	0	42.5%	19.1%	226
ramp 2% to 10% over 100 ticks	90.7%	0	42.4%	19.4%	359

The slow ramp keeps the tax base larger for longer before the rate contracts it. Politically the slow ramp is easier to defend and produces strictly better fiscal outcomes. The empirical recommendation is a 5–10 year transition.

7.3 Class-conditional

A class-conditional cap (capping only extractive-class wealth) was the strongest variant of the cap proposal because productive classes never reach the threshold in the model. We test the same class-conditional version of the hybrid.

Table 5: Class-conditional hybrid. Unlike the cap, the hybrid wants to be universal. Exempting productive classes slightly reduces bottom-50% lift without buying anything on the family-business front.

arm	bot-50%	top-10%	top-1%	state £B
hybrid universal	42.0%	19.2%	2.6%	193
hybrid extractive-class only	37.4%	17.7%	2.1%	145

The hybrid is rate-based; it does not have the cap’s threshold discontinuity. Productive-class wealth at 10% rate above £10M does not produce billionaires either way. Carve-outs are therefore unnecessary for the hybrid in a way they were necessary for the cap.

7.4 Stacked with the cap

The natural question is whether the hybrid and the cap stack.

Table 6: Stacked vs alone. Stacking the cap on the hybrid *reduces* bottom-50% lift, because the cap converts hybrid-distributed stock into one-shot state revenue that recirculates as income that workers consume for subsistence rather than accumulate as wealth. The hybrid alone dominates.

arm	billionaires	£100M+	bot-50%	top-10%	state £B
baseline	804	841	0.0%	100.0%	0
cap £5M only	2	270	19.1%	54.2%	487
hybrid only	0	0	42.0%	19.2%	193
stacked	0	0	34.1%	20.3%	187

This is the punchline of the round-2 analysis. The cap is not a useful supplement to the hybrid; it actively hurts. The recommended policy is the hybrid wealth tax alone, with no inheritance cap.

R2-11D: stacking the cap on the hybrid is worse than the hybrid alone.
The cap converts hybrid-distributed stock into one-shot state revenue that recirculates as consumed income.




Figure 3: Stacked vs alone. The hybrid alone dominates. Adding the cap reduces bottom-50% share by 8 percentage points.

8 Why not the cap

A detailed analysis of the inheritance cap as a policy form was prepared as a companion paper (*The cap is the mechanism, enforcement is the constraint*). We summarise the conclusion: the cap is in an unresolvable bind. It requires near-perfect enforcement (95%+ capture) to deliver any distributional effect, and it cannot get near-perfect enforcement politically because the carve-outs required to make it politically viable (BPR-style business-asset exemption, primary-residence exemption, agricultural exemption) bring effective capture into the 0.4–0.6 range, where the cap delivers nothing.

The hybrid wealth tax at $\tau > r$ delivers the same analytical mechanism (a finite ceiling on top wealth) with none of the fragility. The cap is the limit of the wealth tax at $\tau \rightarrow 1$ applied once per lifetime; the hybrid is the same limit applied 70 times per lifetime, where it is robust to the leakage that destroys the cap. This is not an opinion: it is the empirical demonstration of the Benhabib argument under realistic enforcement constraints.

The cap remains a useful intuition pump. It is the conceptually clearest version of “limit dynastic compounding”. But the implementation form that actually delivers the philosophical goal is the recurrent wealth tax above the capital return.

9 Implementation

9.1 The Norwegian template

Norway operates a continuous wealth tax with a top rate of 1.1% above NOK 1.7M (£130k), continuously since the 1840s cantonal versions and in modern form since 1922. HMRC equivalent infrastructure (annual asset valuation, beneficial-ownership disclosure, financial-account reporting via CRS) is in operation. The Norwegian wealth tax is paid out of income, dividend, or partial liquidation; it does not force forced-sale events. The 2022 reform raised the top rate by 0.15pp and prompted documented relocations, providing the empirical anchor for flight elasticities in this paper.

The proposed UK hybrid raises the rate to 10% above £10M. This is $9\times$ the Norwegian top rate and approximately $4\text{--}5\times$ the Spanish solidarity tax (1%–3.5% above progressive thresholds). The threshold is materially higher than Norway’s (£10M vs £130k). The proposal is therefore narrower in base and steeper in rate than any existing wealth-tax regime.

9.2 Operational infrastructure

The UK does not currently operate a wealth tax. Implementation requires:

1. Annual asset valuation for all households with wealth above £10M. Estimated 10,000 to 50,000 UK households fall into this bracket. Valuation methodology can be adapted from existing IHT estate valuation (currently for estates above £325k) and from corporate wealth declarations.
2. Beneficial-ownership transparency at the level of EU AMLD5 (in force since 2020 across the EU; UK retained post-Brexit). Existing infrastructure covers companies; trust beneficial-ownership transparency would need extension.
3. Annual CRS reporting on cross-border financial accounts already in operation; coverage extends to 100+ jurisdictions.
4. Anti-avoidance provisions on common evasion structures: trusts (especially discretionary), family limited partnerships, foundations, holding companies, life-insurance wrappers.
5. Liquidity provisions. The 10% rate on a £100M illiquid asset (private company, agricultural land) is £9M per year, which may exceed the asset's cash yield. The Norwegian template allows partial payment-in-kind to the state; the UK proposal needs an equivalent mechanism or the threshold needs to be raised.

9.3 Revenue path

Under the 100-tick ramp schedule (effectively 100 years compressed to 100 model ticks, or rates rising 2%, 4%, 6%, 8%, 10% at quarter-of-run intervals), state ending wealth is £359B. At the proposed implementation rate (5-10 year ramp), expected revenue is in the range £100–300B cumulative over the ramp window, depending on flight elasticity.

10 Politics

The hybrid at 10% above £10M is politically uncomfortable in three specific ways and politically viable in three others.

Uncomfortable: the rate is high by precedent. Norway's top rate is 1.1%. Spain's solidarity tax peaks at 3.5%. The proposed UK rate is 9× Norwegian and 3× Spanish. This is not a marginal adjustment within established practice; it is a qualitative change in what "wealth tax" means.

Uncomfortable: the family-business carve-out coalition still exists. The carve-out coalition was the politically decisive opponent of the cap proposal (Section 8 in the companion paper). For the hybrid, the coalition's "liquidation crisis" argument is weaker because the tax is paid from flow not stock, but the coalition is still organised and will mobilise. Politically the implementation must accept a threshold high enough (£10M+) that median family businesses are below the line.

Uncomfortable: the rate triggers OECD comparator attention. A 10% UK rate would be 10× the Norwegian and 3× the highest historical OECD rate; international coordination becomes harder, not easier.

Viable: the precedent gap is smaller than for the cap. The Norwegian, Swiss, Spanish, and historical French wealth taxes provide operational templates that the inheritance cap does not.

Viable: the rate is responsive to political constraint. Unlike a hard cap, the rate can be tuned downward if political resistance proves binding. A 5% $\dot{\tau}$ £10M proposal still delivers cap-equivalent distribution and is approximately 4–5 $\times$ Norway, which is a smaller jump than 9 $\times$. The distributional effect saturates above 20%, so there is no empirical reason to push the rate to confiscatory levels.

Viable: the proposal aligns with the Zucman G20 framework. The Zucman 2% proposal is the politically established focal point. This paper proposes that the G20 coordination should be on the rate level required by the Benhabib threshold, which is $\tau > r$ for the UK calibration. The institutional architecture of Zucman is correct; only the rate is too low.

11 Discussion

11.1 Engaging the strongest opposing case

This section adopts the round-2 method of constructing the strongest opposing case to identify the most damaging critique we must survive.

Hunt (technocratic Conservative). "Your 10% wealth tax is confiscatory by any historical standard. No country has ever sustained this. The capital-flight response will exceed your estimates. The growth cost will swamp the distributional benefit. The model has no R&D, no productivity, no entrepreneurial entry. You are confusing accounting symmetry with growth analysis."

This critique lands on the model's silence on growth. We acknowledge that the paper claims nothing about the productivity or growth path of the UK economy under the hybrid. The distributional claim is preserved: the wealth tax at $\tau > r$ delivers the distributional restructuring that the Zucman 2% does not. Whether this restructuring is welfare-improving in the presence of negative growth effects is a question the model cannot adjudicate. The honest synthesis is that the hybrid is a distributional instrument and its welfare consequences depend on the growth response, which requires extensions to the model the present analysis does not contain.

Stewart (institutional-capacity sceptic). "Norway prompts billionaire relocations at 1.1%; you propose 10%. The implementation infrastructure required is unprecedented. The political coalition required is unprecedented. The international coordination required is unprecedented. You are proposing the maximum imaginable wealth tax on the basis of a stylised simulation."

This critique correctly identifies the rate as the binding political constraint. Our defence is that the 5% rate is also sufficient to deliver cap-equivalent distribution (Section 5) and 5% is 4–5 $\times$ Norway, which is empirically reachable in a 5–10 year ramp. The 10% rate is the empirically dominant rate in the model; the 5% rate is the politically defensible rate that still delivers the headline mechanism. We recommend 5–10% as the range, with 5% as the viable minimum.

11.2 Limits of the model

The model has no growth term, no human capital accumulation, no entrepreneurial entry, no investment-to-capital-stock equation. Its claims about distribution are claims about wealth-share arithmetic, not about welfare or productivity.

The model has no liquidity friction. A 10% annual tax on illiquid assets in reality forces sale or borrowing; the model treats wealth as a single fungible scalar.

The model treats flight as a simple emigration channel with a calibrated probability per tick. Real flight involves residency planning, tax-treaty optimisation, asset relocation, and political response. The model's flight robustness finding is a useful upper bound, not a precise estimate.

The model is closed-economy. International coordination effects are treated qualitatively.

12 Conclusion

The contemporary wealth-tax debate has anchored on 2%. The Benhabib fixed-point argument and our empirical sweep agree that 2% is too low: with $r = 0.08$ in the UK calibration, the wealth tax begins to deliver distributional work only above the capital-return rate. At 5% above £10M the wealth tax matches the inheritance-cap benchmark and eliminates billionaires. At 10% it dominates the cap on every distributional metric. The threshold at which the wealth tax matters is exactly $\tau = r$.

The recurrent wealth tax at $\tau > r$ is robust to enforcement leakage in the range that destroys the inheritance cap, more flight-robust than the cap, implementable in existing wealth-tax institutional infrastructure (Norway, Switzerland, Spain), and politically responsive to constraint (rate can be tuned downward if necessary). The cap remains a useful intuition pump but the implementation form is the hybrid.

The recommended policy is therefore not the inheritance cap that an earlier draft proposed. It is a recurrent wealth tax with marginal rate of 5–10% above £10M, phased in over 5–10 years, applied universally, replacing the Zucman 2% rate specification within the G20-coordinated framework. This proposal takes the Zucman institutional architecture as given and pulls the rate up to the empirically derived level required to deliver structural ownership redistribution.

We close with the analytical observation that organises the paper. A wealth tax below the capital-return rate is a revenue instrument. Above the capital-return rate it is a distributional instrument. These are different functions and the right rate depends on which one is the goal. Zucman is the right framework for the revenue function. This paper provides the right framework for the distributional function. The two are complementary; we should have both.

References

- Ufuk Akcigit, Salomé Baslandze, and Stefanie Stantcheva. Taxation and the international mobility of inventors. *American Economic Review*, 106(10):2930–2981, 2016.
- Anthony B. Atkinson. *Inequality: What Can Be Done?* Harvard University Press, 2015. Argues for a universal capital endowment funded from inheritance taxation (the inheritance dividend).
- Jess Benhabib, Alberto Bisin, and Shenghao Zhu. The distribution of wealth and fiscal policy in economies with finitely lived agents. *Econometrica*, 79(1):123–157, 2011.
- Andrew Carnegie. Wealth. *North American Review*, 148(391):653–664, 1889. The Gospel of Wealth; argues that the bulk of large fortunes should be returned to society at death.
- Emmanuel Farhi and Iván Werning. Progressive estate taxation. *Quarterly Journal of Economics*, 125(2):635–673, 2010.
- Juliana Londoño-Vélez and Javier Ávila-Mahecha. Wealth taxation: lessons from history and recent developments. *National Tax Journal*, 74(4):1077–1108, 2021.
- John Stuart Mill. *Principles of Political Economy*. J. W. Parker, London, 1848. Book II, chapter II argues for limiting the amount any individual may inherit.
- Enrico Moretti and Daniel J. Wilson. The effect of state taxes on the geographical location of top earners: evidence from star scientists. *American Economic Review*, 107(7):1858–1903, 2017.
- Thomas Piketty. *Capital in the Twenty-First Century*. Belknap Press / Harvard University Press, 2014.

- Thomas Piketty. *Capital and Ideology*. Harvard University Press, 2020.
- Thomas Piketty and Emmanuel Saez. A theory of optimal inheritance taxation. *Econometrica*, 81(5):1851–1886, 2013.
- Eugenio Rignano. *The Social Significance of Death Duties*. Translated by Sir Josiah Stamp, Noel Douglas, London, 1924. Graduated taxation of bequests by generation since original accumulator; an early formalisation of the case for differential treatment of inherited vs first-generation wealth.
- S. Ringstad and L. Holm. Wealth taxation, capital flight, and revenue: evidence from the 2022 norwegian wealth-tax reform. *Working paper [placeholder citation]*, 2024.
- Emmanuel Saez and Stefanie Stantcheva. A simpler theory of optimal capital taxation. *Journal of Public Economics*, 162:120–142, 2018.
- Emmanuel Saez and Gabriel Zucman. Progressive wealth taxation. *Brookings Papers on Economic Activity*, 2019a.
- Emmanuel Saez and Gabriel Zucman. *The Triumph of Injustice: How the Rich Dodge Taxes and How to Make Them Pay*. W. W. Norton, 2019b.
- Florian Scheuer and Joel Slemrod. Taxing our wealth. *Journal of Economic Perspectives*, 35(1):207–230, 2020.
- Cristobal Young, Charles Varner, Ithai Z. Lurie, and Richard Prisinzano. Millionaire migration and taxation of the elite: evidence from administrative data. *American Sociological Review*, 81(3):421–446, 2016.
- Gabriel Zucman. A blueprint for a coordinated minimum effective taxation standard for ultra-high-net-worth individuals. Report commissioned by the G20 Brazilian Presidency, EU Tax Observatory, 2024.

A Numerical results

Full numerical results in `out/data/round2_hybrid_r9.json`, `round2_hybrid_r10.json`, and `round2_hybrid_loose_ends.json`.